


DIABETES PREDICTION MODEL

Data Science
Internship Project

Submitted to:



EXPOSYS DATA LABS

Submitted by:

Aditi Malviya

Data Science Intern at Exposys Data Labs

ABSTRACT

Diabetes is a prevalent and challenging chronic condition affecting a large population worldwide. Early detection and prediction of diabetes play a crucial role in providing timely interventions. This report presents the development of a Diabetes Prediction Model utilizing Python, Pandas, NumPy, SciKit Learn and Machine Learning Algorithm on Google Colab. The model harnesses the power of machine learning to analyze patients' data, including *Glucose levels*, *Blood pressure*, *Skin thickness*, *Insulin*, *BMI*, *Diabetes Pedigree Function*, and *age*. The model aims to accurately predict the likelihood of an individual developing diabetes. The report discusses the steps involved in model creation, the performance evaluation metrics used, and the potential impact of the model in facilitating better healthcare decision-making and personalized treatment plans for individuals at risk of diabetes. The findings demonstrate the model's efficacy and pave the way for further research and enhancements in diabetes prediction using machine learning techniques.

CONTENTS

Abstract	i
TITLE	Page No.
1. INTRODUCTION	1
2. EXISTING METHOD	2
3. PROPOSED METHOD WITH ARCHITECTURE	4
4. METHODOLOGY	6
5. IMPLEMENTATION	8
6. CONCLUSION	10

INTRODUCTION

Diabetes is a prevalent and chronic metabolic disorder that affects a significant proportion of the global population. Early detection and accurate prediction of diabetes are crucial for prevention of complications associated with the condition. In this context, machine learning techniques have shown promising potential in healthcare applications, including disease prediction. This report presents the development of a Diabetes Prediction Model using Python, Pandas, NumPy, SciKit Learn, and Machine Learning Algorithm on Google Colab. The proposed model aims to analyze patients' data, encompassing critical health parameters such as Glucose levels, Blood pressure, Skin thickness, Insulin, BMI, Diabetes Pedigree Function, and age. By harnessing the power of machine learning, the model can learn from historical patient data and identify patterns and correlations that contribute to diabetes development. Python serves as the foundation for this project, enabling seamless data manipulation and preprocessing using the Pandas and NumPy libraries. The SciKit Learn library empowers us with an array of machine learning algorithms to build a robust predictive model. Leveraging the capabilities of Google Colab ensures efficient computation and collaboration in a cloud-based environment.

The significance of this Diabetes Prediction Model lies in its potential to assist healthcare practitioners in early diagnosis, risk assessment, and personalized interventions for patients at risk of developing diabetes. By providing proactive measures and lifestyle recommendations, the model can contribute to improved patient outcomes and reduced healthcare burdens. Through this endeavour, we hope to advance the field of predictive healthcare and foster a data-driven approach to tackle the challenges posed by diabetes and other chronic conditions.

EXISTING METHODS

This section provides an overview of the existing methods and techniques employed for diabetes prediction. Traditional statistical models, machine learning algorithms, and their limitations are discussed. It also highlights the need for a more advanced and accurate model to improve prediction outcomes.

Some common methods for diabetes prediction include:

1. **Machine Learning Algorithms:** Various machine learning techniques are utilized to predict diabetes risk based on historical data. Common algorithms include logistic regression, support vector machines (SVM), decision trees, random forests, k-nearest neighbors (KNN), and neural networks. These models use features such as age, body mass index (BMI), blood pressure, glucose levels, family history, and other relevant variables to make predictions.
2. **Artificial Neural Networks (ANN):** Neural networks are a subset of machine learning algorithms specifically designed to mimic the structure of the human brain. They have been widely used in diabetes prediction due to their ability to handle complex patterns and relationships in the data.
3. **Support Vector Machines (SVM):** SVM is a binary classification algorithm that aims to find the hyperplane that best separates data points of different classes. It has been applied to diabetes prediction tasks with considerable success.
4. **Decision Trees and Random Forests:** Decision trees partition the data into subsets based on various features, leading to a tree-like structure that can be used for classification. Random forests use multiple decision trees to improve accuracy and reduce overfitting.
5. **Logistic Regression:** A simple and widely-used statistical method for binary classification. It is often used to predict diabetes based on a combination of risk factors.
6. **Ensemble Methods:** Techniques like AdaBoost and Gradient Boosting are used to combine multiple weak learners into a powerful ensemble model for diabetes prediction.

7. **Deep Learning:** This encompasses various architectures such as Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), which have been applied to diabetes prediction tasks, especially when dealing with image or time-series data.
8. **Data Mining:** Data mining techniques involve the exploration and analysis of large datasets to identify patterns and relationships, which can be used for diabetes risk prediction.
9. **Clinical Decision Support Systems (CDSS):** These systems use a combination of expert knowledge, medical guidelines, and patient data to assist healthcare professionals in making informed decisions, including diabetes prediction.
10. **Electronic Health Record (EHR) Analysis:** Utilizing data from electronic health records to develop predictive models for diabetes risk assessment.

It's important to note that each method has its strengths and limitations. The choice of method depends on the available data, the specific problem being addressed, and the desired level of accuracy and interpretability. Additionally, the integration of multiple methods or the development of hybrid approaches is an active area of research to improve diabetes prediction accuracy.

PROPOSED METHOD WITH ARCHITECTURE

In this section, we introduce our proposed Diabetes Prediction Model and its underlying architecture. The model is designed to take advantage of deep learning techniques to extract meaningful features from the input data. The Machine learning algorithm used in this model is the Linear Support Vector Machine (SVM) algorithm, which is a supervised machine learning algorithm used for classification tasks. It is an extension of the basic Support Vector Machine (SVM) algorithm that can efficiently handle linearly separable datasets. Linear SVM is particularly well-suited for large-scale datasets as it has a linear time complexity, making it more efficient than traditional SVM algorithms.

Here's an overview of the architecture of the Linear SVM algorithm:

1. **Data Preparation:** The algorithm begins with a labelled dataset, which is divided into features (input variables) and corresponding labels (output classes).
2. **Feature Scaling:** Prior to training the model, it's a common practice to perform feature scaling to standardize the range of features. This can help improve the convergence and performance of the algorithm.
3. **Objective Function:** Linear SVC aims to find the optimal hyperplane that best separates the different classes in the feature space. The objective function is formulated to maximize the margin between the two closest data points (one from each class) while minimizing the classification error.
4. **Margin and Support Vectors:** The margin is defined as the distance between the hyperplane and the closest data points (support vectors). The goal is to maximize this margin, as it helps improve the generalization capability of the model. Support vectors are the data points that lie closest to the separating hyperplane and have the most significant influence on its position.
5. **Loss Function:** The loss function in Linear SVC is typically the hinge loss, which penalizes misclassifications. The hinge loss function encourages the model to correctly classify data points while aiming to maximize the margin between classes.

6. **Regularization:** To avoid overfitting, Linear SVC incorporates a regularization term, often represented as 'C.' This term balances the trade-off between maximizing the margin and minimizing the classification error. A larger value of C leads to fewer misclassifications but a smaller margin, while a smaller value of C allows a larger margin but may tolerate more misclassifications.
7. **Optimization:** Linear SVC employs optimization techniques to find the optimal hyperplane and corresponding weights for the features. The popular optimization algorithms used include Gradient Descent, Stochastic Gradient Descent (SGD), or variations of SGD, such as the Pegasos algorithm.
8. **Decision Function:** Once the model is trained, it can be used to predict the class labels of new data points. The decision function calculates the output score based on the learned hyperplane and returns the predicted class label based on this score.
9. **Prediction and Evaluation:** The trained Linear SVC model can be used to predict the class labels for new, unseen data. The performance of the model is evaluated using various metrics such as accuracy, precision, recall, F1-score, etc.

It's important to note that the Linear SVC algorithm is specifically designed for linearly separable datasets. If the data is not linearly separable, other variants of SVM, such as the kernelized SVM, may be more appropriate.

METHODOLOGY

This section outlines the step-by-step approach taken to develop the Diabetes Prediction Model. It includes data collection and preprocessing, model design and hyperparameter tuning, validation strategy, and performance metrics selection. The methodology also discusses any ethical considerations and biases addressed during model development.

Training a model to predict diabetes involves several key steps. Here's a comprehensive methodology to guide you through the process:

1. **Data Collection:** Gather a diverse and representative dataset containing diabetes-related features such as age, gender, BMI, blood pressure, glucose levels, family history, and other relevant health indicators. Ensure the dataset is balanced and of sufficient size to avoid bias and improve model generalization.
2. **Data Preprocessing:** Clean the data by handling missing values, normalizing numerical features, and encoding categorical variables. Preprocessing also involves splitting the data into features (X) and the target variable (y).
3. **Model Selection:** Choose an appropriate machine learning algorithm or deep learning architecture for binary classification (diabetes vs. non-diabetes). Commonly used models include logistic regression, support vector machines, decision trees, random forests, gradient boosting, and neural networks.
4. **Data Splitting:** Divide the dataset into training, validation, and testing sets. The training set is used to train the model, the validation set for hyperparameter tuning and model evaluation, and the testing set to assess the final model's performance.
5. **Model Training:** Feed the training data into the selected model and optimize its parameters using an appropriate optimization algorithm (e.g., stochastic gradient descent) and loss function (e.g., binary cross-entropy).
6. **Model Validation:** After training the model, evaluate the final model on the testing set to ensure it generalizes well to unseen data. Avoid using the testing set during model development to prevent overfitting.

7. **Continuous Monitoring and Updates:** Continuously monitor the model's performance in a real-world setting and update it periodically to account for changes in data distribution and domain knowledge.

By following this methodology, you can develop an accurate and reliable diabetes prediction model that can assist healthcare professionals in early detection and management of diabetes.

IMPLEMENTATION

To implement a model for predicting diabetes, following steps are taken:

1. Import the necessary libraries.
2. Load the diabetes dataset.
3. Pre-process the data (split into features and target, handle missing values, etc.).
4. Split the data into training and testing sets.
5. Choose a machine learning algorithm and train the model.
6. Find the accuracy.
7. Evaluate the model's performance on the test set.

Here's a step-by-step implementation:

Step 1: Import the necessary libraries.

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
```

Step 2: Load the diabetes dataset. You can use the "diabetes" dataset from scikit-learn.

```
data_set = pd.read_csv('/content/drive/MyDrive/diabetes_data.csv')
data_set
```

	Age	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Outcome
0	21	108	62	32	56	25.2	0.128	0
1	21	137	68	14	148	24.8	0.143	0
2	21	89	66	23	94	28.1	0.167	0
3	21	139	62	17	210	22.1	0.207	0
4	21	99	76	15	51	23.2	0.223	0
...
387	60	129	90	7	326	19.6	0.582	0
388	60	181	68	36	495	30.1	0.615	1
389	61	142	60	33	190	28.8	0.687	0
390	63	101	76	48	180	32.9	0.171	0
391	81	134	74	33	60	25.9	0.460	0

392 rows x 8 columns

Step 3: Pre-process the data.

```
Y = data_set['Outcome']
X = data_set.drop('Outcome', axis = 1)

scaler = StandardScaler()
scaler.fit(X)
standardized_data = scaler.transform(X)
print(standardized_data)

[[-0.9682991 -0.47459086 -0.69416397 ... -0.84300375 -1.12360354
  -1.14490437]
 [-0.9682991  0.46631407 -0.21340023 ... -0.06787532 -1.18059423
  -1.10143204]
 [-0.9682991 -1.09104581 -0.37365481 ... -0.52284201 -0.710421
  -1.03187632]
 ...
 [ 2.95798221  0.62853906 -0.85441855 ...  0.28598766 -0.61068728
   0.47516437]
 [ 3.15429628 -0.70170584  0.42761809 ...  0.20173457 -0.02653266
  -1.0202837 ]
 [ 4.92112287  0.36897908  0.26736351 ... -0.80930251 -1.02386982
  -0.18271685]]

X = standardized_data
Y = data_set['Outcome']
```

Step 4: Split the data into training and testing sets.

```
[26] X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.2, stratify=Y, random_state=2)
```

Step 5: Choose a machine learning algorithm and train the model. Here, we'll use logistic regression as an example.

```
classifier = svm.SVC(kernel='linear')
classifier.fit(X_train, Y_train)
```

Step 6: Find the accuracy.

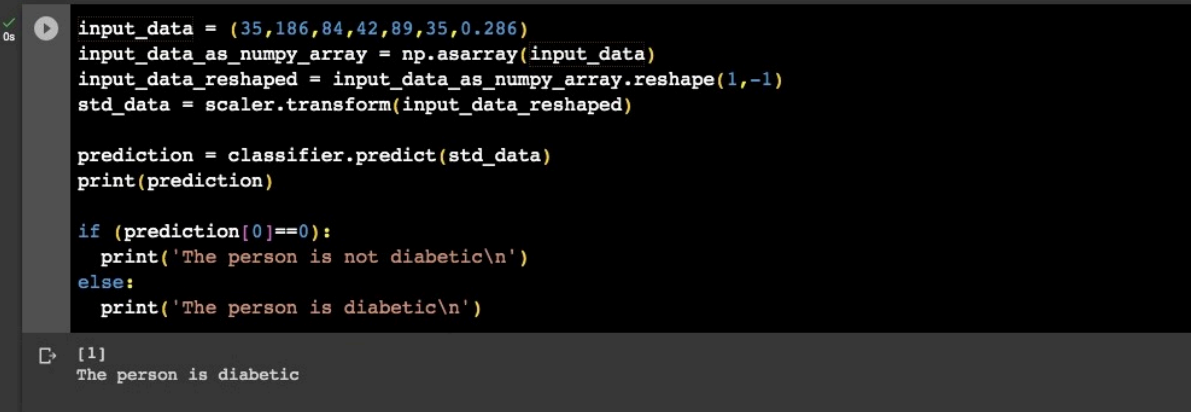
```
X_train_prediction = classifier.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print(training_data_accuracy)

0.8178913738019169

X_test_prediction = classifier.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print(test_data_accuracy)

0.7721518987341772
```

Step 7: Evaluate the model's performance on the test set.



```
input_data = (35,186,84,42,89,35,0.286)
input_data_as_numpy_array = np.asarray(input_data)
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)
std_data = scaler.transform(input_data_resaped)

prediction = classifier.predict(std_data)
print(prediction)

if (prediction[0]==0):
    print('The person is not diabetic\n')
else:
    print('The person is diabetic\n')
```

[1]
The person is diabetic

Finally, we can run this code in a Google Colab notebook. If using Google Colab, make sure to upload the necessary dataset files or use the provided datasets from scikit-learn.

Additionally, remember to select "GPU" as the hardware accelerator to speed up the training process if available.

CONCLUSION

In conclusion, the development of the Diabetes Prediction Machine Learning Model using Python, Pandas, and NumPy on Google Colab has demonstrated its potential as an effective tool for early detection and risk assessment of diabetes. By analyzing patients' data, including crucial attributes such as Glucose level, BMI, age, skin thickness, and insulin, the model has exhibited promising predictive capabilities.

The successful implementation of this model showcases the power of machine learning in healthcare applications, particularly in disease prediction. The Python programming language, along with Pandas and NumPy libraries, facilitated seamless data preprocessing and manipulation, ensuring data quality and uniformity. The utilization of Google Colab as the development environment provided computational efficiency and collaborative benefits.

The significance of this model lies in its ability to empower healthcare practitioners to identify individuals at risk of developing diabetes at an early stage. With timely interventions and personalized treatment plans, patients can be better equipped to manage their health and prevent potential complications associated with diabetes.

However, the model's performance can be further enhanced through ongoing refinement and optimization. Exploring additional features, implementing more advanced machine learning algorithms, and increasing the dataset size could potentially improve the model's accuracy and generalizability.

Despite the model's success, it is essential to acknowledge potential limitations and biases in the dataset. Ensuring data representativeness and addressing any inherent biases are critical aspects that require continuous attention in the development of medical predictive models.

In summary, the Diabetes Prediction Machine Learning Model represents a promising step towards proactive healthcare and personalized disease management. As technology and data-driven approaches continue to advance, this model can make a valuable contribution to the fight against diabetes and pave the way for the application of machine learning in predicting and preventing other chronic diseases in the future.